

Assignment #1

Ivana Nworah Bortot and Irene Avezzi


1. Similarity and Distance Measures Between Data

Do the same of exercise 1 in Lab 2 with the digits dataset. This dataset can be downloaded and inspected with the Sklearn APIs. However, it has a different format with respect to the `train_catvnoncat.h5` dataset.

Try to understand the data format by using the following code:

```
1 from sklearn.datasets import load_digits
2 import matplotlib.pyplot as plt
3
4 digits = load_digits()
5 print(digits.data.shape)
6
7 plt.gray()
8
9 plt.imshow(digits.images[0])
```

(1797, 64)



You can use the code written for the `L1`, `L2` and `cosine` metrics. However, the code for the `top1Similarity` has to be slightly modified. In this new version, the flattened version of the dataset is a matrix where each row is an image index and the columns are the pixel values. Notice that digit images have only one channel.

Answer the following question by properly coding, testing and discussing:

- for each digit from 0 to 9, are the `L1`, `L2` and `cosine` similarity able to return similar images? That is, an image containing the same digit.

```
1 import numpy as np
```

```
1 print(f"images shape: {digits.images.shape}")
2
3 dataset_flatten = digits.images.reshape(-1, digits.images.shape[0]).T
4
5 print(f"images flatten shape: {dataset_flatten.shape}")
```

images shape: (1797, 8, 8)
images flatten shape: (1797, 64)

```

1 #L1 Metric
2
3 from math import *
4
5 def L1(x1, x2):
6     """
7     Arguments:
8     x1 -- vector of size m
9     x2 -- vector of size m
10
11     Returns:
12     distance -- the L1 distance between the two vectors
13     """
14
15     x1 = x1.astype(np.int64)
16     x2 = x2.astype(np.int64)
17
18     distanceL1 = np.sum(np.abs(x2 - x1))
19     return distanceL1

```

```

1 # L2 Metric
2
3 def L2(x1, x2):
4     """
5     Arguments:
6     x1 -- vector of size m
7     x2 -- vector of size m
8
9     Returns:
10    distance -- the L2 distance between the two vectors
11    """
12
13    x1 = x1.astype(np.int64)
14    x2 = x2.astype(np.int64)
15    distanceL2 = np.sqrt(np.sum(np.square(x2 - x1)))
16    return distanceL2

```

```

1 # Cosine Similarity
2
3 def Cosine(x1, x2):
4     """
5     Arguments:
6     x1 -- vector of size m
7     x2 -- vector of size m
8
9     Returns:
10    distance -- the cosine distance between the two vectors
11    """
12
13    x1 = x1.astype(np.int64)
14    x2 = x2.astype(np.int64)
15    simCos = np.sum(x1 * x2)/(np.sqrt(np.sum(x1 * x1)) * np.sqrt(np.sum(x2 * x2)))
16    distanceCos = 1 - simCos #the cosine similarity must be converted into a distance
17
18    return distanceCos

```

```

1 def top1Similarity(img, dataset, checkReturnQuery): #checkReturnQuery=true include query, false=non include query
2     """
3     Arguments:
4     img -- vector of size m (represents an image)
5     dataset -- a DataFrame of images where each column represents an image (a vector of size m)
6
7     Returns:
8     top1 -- the most similar image
9     """
10
11     # Compute L1, L2, and Cosine Distances
12     distanceL1 = L1(dataset[:,img], dataset[:,0])
13     distanceL2 = L2(dataset[:,img], dataset[:,0])
14     distanceCos = Cosine(dataset[:,img], dataset[:,0])
15
16     topL1_Image = 0
17     topL2_Image = 0
18     topCos_Image = 0
19
20     if (checkReturnQuery):
21         limit = -1 #always true
22     else:
23         limit = img # skip the input img
24
25     for i in range(1, dataset.shape[1]):
26         if i != limit:
27             L1value = L1(dataset[:,img], dataset[:,i])
28             L2value = L2(dataset[:,img], dataset[:,i])
29             Cosvalue = Cosine(dataset[:,img], dataset[:,i])
30
31             if L1value < distanceL1:# or distanceL1 == 0:
32                 distanceL1 = L1value
33                 topL1_Image = i
34
35             if L2value < distanceL2:# or distanceL2 == 0:
36                 distanceL2 = L2value
37                 topL2_Image = i
38
39             if Cosvalue < distanceCos:# or distanceCos == 0:
40                 distanceCos = Cosvalue
41                 topCos_Image = i
42
43
44     return topL1_Image, topL2_Image, topCos_Image, distanceL1, distanceL2, distanceCos

```


```

1 # Test the Similarity Function
2 QueryImage = 2
3 topL1_Image, topL2_Image, topCos_Image, distanceL1, distanceL2, distanceCos = top1Similarity(QueryImage, dataset_flatten, False)
4 print("The query image is:")
5 plt.imshow(digits.images[QueryImage])
6 print("The query id is: ", QueryImage, "; the L1 image id is: ", topL1_Image, "; the L2 image id is: ", topL2_Image, ";

```


The query image is:

The query id is: 2 ; the L1 image id is: 50 ; the L2 image id is: 50 ; the cos image id is: 26




```
1 print(f"The most similar image using L1 is (score {distanceL1}):")
2 plt.imshow(digits.images[topL1_Image])
```

The most similar image using L1 is (score 7148):




```
1 print(f"The most similar image using L2 is (score {distanceL2}):")
2 plt.imshow(digits.images[topL2_Image])
```

The most similar image using L2 is (score 264.16661409042587):



```
1 print(f"The most similar image using Cosine is (score {distanceCos}):")
2 plt.imshow(digits.images[topCos_Image])
```

The most similar image using Cosine is (score 0.32061224028095103):



2. Data Preprocessing

This exercise requires the understanding and pre-processing of the Car Evaluation dataset. You can find the dataset and its description here <https://archive-beta.ics.uci.edu/dataset/19/car+evaluation>.

The tasks to perform are the following:

1. Download the dataset, understand the data and load it with Pandas as a data frame. Pay attention that you also need to set the name of the columns.
2. How many attributes are in the dataset? What are the types of these attributes?
3. Check the presence of None values.
4. The attributes are encoded as string, transform these ordinal attributes into numbers.
5. Show the data distribution, that is, the histogram of each column. Is there anything relevant to discuss for each attribute?
6. Show the boxplots of the attributes. Discuss the presence of outliers.
7. Remove the outliers (if present) for the numeric attributes by using the IQR method.
8. Select the data about the acceptable and unacceptable cars. Is there imbalanced in the classes? You can compute such a balance by dividing the number of acceptable cars by the total number of samples.
9. Plot the scatter plot divided per classes.
10. Are the classes easily separable?
11. By inspecting the scatter plots, are their important attributes that allow an easy separation of the classes?
12. Are there correlated features?

1. Download the dataset, understand the data and load it with Pandas as a dataframe.

```
1 import pandas as pd
2 data = pd.read_csv('car.data.csv', header=None)
3 data.columns = ['buying', 'maint', 'doors', 'persons', 'lug_boot', 'safety', 'class']
```

2. How many attributes are in the dataset? What are the types of these attributes?

```
1 print(f'Number of attributes = {data.shape[1]}')
2 data.info()
```

```
Number of attributes = 7
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1728 entries, 0 to 1727
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   buying      1728 non-null   object
1   maint       1728 non-null   object
2   doors       1728 non-null   object
3   persons     1728 non-null   object
4   lug_boot    1728 non-null   object
5   safety      1728 non-null   object
6   class       1728 non-null   object
dtypes: object(7)
memory usage: 94.6+ KB
```

3. Check the presence of none values

```
1 import numpy as np
2
3 data = data.replace('?', np.NaN)
4
5 print('Number of missing values:')
6 for col in data.columns:
7     print(f'\t{col}: {data[col].isna().sum()}')
```

```
Number of missing values:
buying: 0
maint: 0
doors: 0
persons: 0
lug_boot: 0
safety: 0
class: 0
```

Si poteva già dedurre che non ci fossero null values osservando lo step 2 perché nelle info del dataset ogni attributo ha tanti valori non-null tante quante sono le istanze del dataset

4. The attributes are encoded as string, transform these ordinal attributes into numbers.

```
1 # Converting Ordinal Variable buying to numeric
2 data['buying'].replace({'vhigh':4, 'high':3, 'med':2, 'low':1}, inplace=True)
3 data['maint'].replace({'vhigh':4, 'high':3, 'med':2, 'low':1}, inplace=True)
4 data['doors'].replace({'5more':5, '4':4, '3':3, '2':2}, inplace=True)
5 data['persons'].replace({'more':5, '4':4, '2':2}, inplace=True)
6 data['lug_boot'].replace({'big':3, 'med':2, 'small':1}, inplace=True)
7 data['safety'].replace({'high':3, 'med':2, 'low':1}, inplace=True)
8 data['class'].replace({'vgood':4, 'good':3, 'acc':2, 'unacc':1}, inplace=True)
9
10 print(data)
11 data.info()
```

	buying	maint	doors	persons	lug_boot	safety	class
0	4	4	2	2	1	1	1
1	4	4	2	2	1	2	1
2	4	4	2	2	1	3	1
3	4	4	2	2	2	1	1
4	4	4	2	2	2	2	1
...
1723	1	1	5	5	2	2	3
1724	1	1	5	5	2	3	4
1725	1	1	5	5	3	1	1
1726	1	1	5	5	3	2	3
1727	1	1	5	5	3	3	4

Dopo la conversione tutti gli attributi sono di tipo int64 e non più object


```
[1728 rows x 7 columns]
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1728 entries, 0 to 1727
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   buying      1728 non-null   int64
1   maint       1728 non-null   int64
2   doors       1728 non-null   int64
3   persons     1728 non-null   int64
4   lug_boot    1728 non-null   int64
5   safety      1728 non-null   int64
6   class       1728 non-null   int64
dtypes: int64(7)
memory usage: 94.6 KB
```

5. Show the data distribution, that is, the histogram of each column. Is there anything relevant to discuss for each attribute?

```
1 #import matplotlib library
2 import matplotlib.pyplot as plt
```


```
1 # plot the histogram of the buying column
2 data['buying'].plot(kind="hist")
3 plt.ylabel('buying')
```

Text(0, 0.5, 'buying')




```
1 # plot the histogram of the maint column
2 data['maint'].plot(kind="hist")
3 plt.ylabel('maint')
```

Text(0, 0.5, 'maint')




```
1 # plot the histogram of the doors column
2 data['doors'].plot(kind="hist")
3 plt.ylabel('doors')
```

Text(0, 0.5, 'doors')




```
1 # plot the histogram of the persons column
2 data['persons'].plot(kind="hist")
3 plt.ylabel('persons')
```

Text(0, 0.5, 'persons')





```
1 # plot the histogram of the lug_boot column
2 data['lug_boot'].plot(kind="hist")
3 plt.ylabel('lug_boot')
```

Text(0, 0.5, 'lug_boot')



```
1 # plot the histogram of the safety column
2 data['safety'].plot(kind="hist")
3 plt.ylabel('safety')
```

Text(0, 0.5, 'safety')



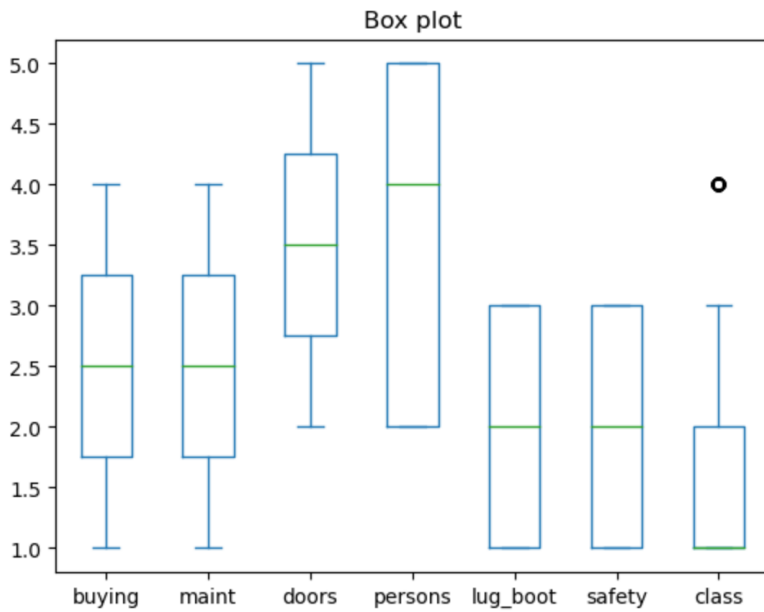
Q: Is there anything relevant to discuss for each attribute?

A: Tutti gli attributi sono equamente distribuiti per i possibili valori

6. Show the boxplots of the attributes. Discuss the presence of outliers.

```
1 data[['buying', 'maint', 'doors', 'persons', 'lug_boot', 'safety', 'class']].plot(kind='box', title='Box plot')
```

<Axes: title={'center': 'Box plot'}>



Solo l'attributo 'class' presenta outliers. Infatti solo una minoranza di istanze rientra nella categoria 'vgood' (sostituita col valore cardinale 4).

7. Remove the outliers (if present) for the numeric attributes by using the IQR method.

```

1 clean_data = data[['class']].copy() #class is the only attribute with outliers
2 print(clean_data)
3
4 Q1 = clean_data.quantile(0.25)
5 Q3 = clean_data.quantile(0.75)
6 IQR = Q3 - Q1
7 print('IQR is: ', IQR)
8
9 clean_data_no_outliers = clean_data.loc[((Q1 - 1.5*IQR < clean_data).sum(axis=1)==len(clean_data.columns)) & ((clean_data
10 & ((clean_data < Q3 + 1.5*IQR).sum(axis=1)==len(clean_data.columns)), :)]
11 print(f"Num of samples before outlier cleaning: {len(clean_data)}")
12 print(f"Num of samples after outlier cleaning: {len(clean_data_no_outliers)}")
13 clean_data_no_outliers.boxplot(figsize=(10,5))
14

```

```


class
0      1
1      1
2      1
3      1
4      1
...    ...
1723   3
1724   4
1725   1
1726   3
1727   4

```

```

[1728 rows x 1 columns]
IQR is: class    1.0
dtype: float64
Num of samples before outlier cleaning: 1728
Num of samples after outlier cleaning: 1663

```



Abbiamo eseguito la rimozione degli outliers solamente per l'attributo class perchè solamente questo attributo presenta outliers

8. Select the data about the acceptable and unacceptable cars. Is there imbalanced in the classes? You can compute such a balance by dividing the number of acceptable cars by the total number of samples.

```

1 acc_data = data[data['class'] == 2] + data[data['class'] == 3] + data[data['class'] == 4]
2 unacc_data = data[data['class'] == 1]
3
4 print(f"Num of acc car: {len(acc_data)}")
5 print(f"Num of unacc car: {len(unacc_data)}")
6 answer = str(round(len(acc_data)/len(data), 6))
7 print(f"Ratio between acc and total number of sample: {answer}")
8
9

```

```

Num of acc car: 518
Num of unacc car: 1210
Ratio between acc and total number of sample: 0.299769

```


Note: abbiamo considerato anche i valori good e vgood nella classe dei dati accurate

Q: Is there imbalanced in the classes?

A: Osserviamo un disequilibrio tra le due classi di dati. Per risolvere questo problema possiamo eseguire oversampling sulla classe di 'acc car' oppure undersampling sulla classe di 'unacc car'

9. Plot the scatter plot divided per classes.

```
1 import matplotlib.pyplot as plt
2
3 index = 0 # number of cell we are going to fill, from 0 to (num_attr^2 - num_attr)/2 = 21
4
5 # the number of rows x number of cols must results in the total number of cells
6 num_rows = 7
7 num_cols = 3
8
9 #this is index of each pair of attributes and goes from 0 to 7
10 num_attrs = len(data.columns)
11
12 fig, axes = plt.subplots(num_rows, num_cols, figsize=(20,30))
13
14 for i in range(num_attrs):
15     for j in range(i + 1, num_attrs):
16         # ax1 represent the col of axes that we are going to fill and ax2 the row
17         ax1 = int(index/num_cols) # from 0 to 6
18         ax2 = index % num_cols # either 0, 1 or 2
19         for class_name in data['class'].unique():
20             # boolean mask for selecting samples of a given class
21             class_df = data[data['class'] == class_name]
22
23             # scatter plot for each couple (i, j) of features
24             axes[ax1][ax2].scatter(class_df.iloc[:, i], class_df.iloc[:, j])
25
26         axes[ax1][ax2].set_xlabel(data.columns[i])
27         axes[ax1][ax2].set_ylabel(data.columns[j])
28         index = index + 1
```



10. Are the classes easily separable?

Generalmente le classi sono facilmente separabili perché raggruppabili chiaramente in base al valore dell'attributo sulle ascisse o sulle ordinate.

11. By inspecting the scatter plots, are there important attributes that allow an easy separation of the classes?

Tutte le coppie di attributi permettono di separare le classi.

Generalmente si viene a creare una separazione lineare tranne nei caso di:

- buying e maint
- maint e lug_boot
- maint e safety
- person e safety
- lug_boot e safety

Solamente alcune coppie di attributi permettono di individuare e separare 4 classi:

- buying e maint
- buying e class
- maint e persons
- maint e safety
- maint e class
- doors e class
- persons e class
- lug_boot e safety
- lug_boot e class
- safety e class

Negli altri casi invece si possono individuare:

- 1 classe:
 - doors e lug_boot 1
- 2 classi:
 - buying e doors
 - buying e lug_boot
 - doors e persons
 - doors e safety
 - persons e lug_boot
 - persons e safety
- 3 classi:
 - buying e persons
 - buying e safety
 - maint e doors
 - maint e lug_boot

12. Are there correlated features?

Non si osserva nessuna correlazione tra le features